

Matrix: Multihop Address allocation and dynamic any-To-any Routing for 6LoWPAN

Bruna S. Peres
bperes@dcc.ufmg.br

Otávio A. de O. Souza
oaugusto@dcc.ufmg.br

Bruno P. Santos
bruno.ps@dcc.ufmg.br

Edson R. A. Junior
edsonroteia@dcc.ufmg.br

Olga Goussevskaya
olga@dcc.ufmg.br

Marcos A. M. Vieira
mmvieira@dcc.ufmg.br

Luiz F. M. Vieira
lfvieira@dcc.ufmg.br

Antonio A.F. Loureiro
loureiro@dcc.ufmg.br

Computer Science Department, Universidade Federal de Minas Gerais, Brazil

ABSTRACT

Standard routing protocols for IPv6 over Low power Wireless Personal Area Networks (6LoWPAN) are mainly designed for data collection applications and work by establishing a tree-based network topology, which enables packets to be sent upwards, from the leaves to the root, adapting to dynamics of low-power communication links. The routing tables in such unidirectional networks are very simple and small, since each node just needs to maintain the address of its parent in the tree, providing the best-quality route at every moment. In this work we propose Matrix, a platform-independent routing protocol that utilizes the existing tree structure of the network to enable reliable and efficient any-to-any data traffic. Matrix uses hierarchical IPv6 address assignment in order to optimize routing table size, while preserving bidirectional routing. Moreover, it uses a local-broadcast mechanism to forward messages to the right subtree when persistent node or link failures occur. We implemented Matrix on TinyOS and evaluated its performance both analytically and through simulations on TOSSIM. Our results showed that the proposed protocol is superior to available protocols, such as RPL and XCTP, when it comes to any-to-any data communication, in terms of reliability, message efficiency and memory footprint.

CCS Concepts

•Networks → Network protocol design; Network layer protocols;

Keywords

6LoWPAN; IPv6; CTP; RPL; any-to-any routing; fault tolerance

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

1. INTRODUCTION

IPv6 over Low-power Wireless Personal Area Networks (6LoWPAN¹) is a working group inspired by the idea that even the smallest low-power devices should be able to run the Internet Protocol to become part of the Internet of Things. Standard routing protocols for 6LoWPAN, such as CTP (Collection Tree Protocol [6]) and RPL (IPv6 Routing Protocol for Low-Power and Lossy Networks [19]), have two distinctive characteristics: communication devices use unstructured IPv6 addresses that do not reflect the topology of the network (typically derived from their MAC addresses), and routing is based on distributed collection tree structures, focused on bottom-up data flows (from the leaves to the root) and lacking support for any-to-any communication. The problem with such one-directional routing is it makes it inefficient to build important network functions, such as configuration routines and reliable mechanisms to ensure the delivery of data end-to-end.

Even though CTP does not support any-to-any traffic, the specification of RPL defines two modes of operation for top-down data flows: the non-storing mode, which uses source routing, and the storing mode, in which each node maintains a routing table for all possible destinations. This requires $O(n)$ space (where n is the total number of nodes), which is unfeasible for memory-constrained devices. Our experiments show that in random topologies with one hundred nodes, with no link or node failures, RPL succeeds to deliver less than 20% of top-down messages sent by the root (TODO REF RPL FIG).

Some works have addressed this problem from different perspectives. In [16], the authors propose a hierarchical IPv6 address allocation scheme, referred to as MHCL, to enable any-to-any routing by incorporating network topology information into the IPv6 address assigned to each node in a multihop fashion. MHCL was implemented as a subroutine of RPL and showed to enable any-to-any routing using compact routing tables. MHCL stores the IPv6 address range of the entire subtree rooted at a child-node in a single routing table entry, which results in $O(k)$ memory space, where k is the number of one-hop descendants of a node in the collection tree. The downside of MHCL is it was not designed to

¹We use the acronym 6LoWPAN to refer to Low power Wireless Personal Area Networks that use IPv6

deal with network faults and dynamic network topologies, since a message can be dropped whenever a persistent node or link failure occurs in the address hierarchy.

In this work, we build upon the idea of using hierarchical IPv6 address allocation and propose Matrix, a routing scheme for dynamic network topologies and fault-tolerant any-to-any data flows in 6LoWPAN. Matrix assumes there is an underlying collection tree topology (provided by CTP or RPL, for instance), in which nodes have static locations, i.e., are not mobile, and links are dynamic, i.e., nodes might choose different parents according to link quality dynamics. Matrix uses only one-hop information in the routing tables and implements a local broadcast mechanism to forward messages to the right subtree when (persistent) node or link failures occur.

After the network has been initialized and all nodes have received an IPv6 address range, three simultaneous distributed trees are maintained by all nodes: the collection tree (Ctree), the IPv6 address tree (IPtree), and the reverse collection tree (RCtree). Initially, IPtree has the same topology as Ctree, RCtree has no links, and any-to-any packet forwarding is performed using Ctree for bottom-up and IPtree for top-down data flows, respectively. Whenever a change occurs in a link of Ctree, the new link is added in reverse direction into RCtree and maintained as long as this topology change persists. Therefore RCtree is not really a tree, since it contains only the reversed links present in Ctree but not in IPtree. When a node or link fails or changes in Ctree, RCtree reflects this change, and packets are forwarded from IPtree to RCtree via a local broadcast. The node that receives a local-broadcast checks in its RCtree table whether it knows the subtree of the destination IPv6 address: if yes, then it forwards the packet to the right subtree via RCtree, and the packet continues its path in the IPtree until the final destination. TODO: RESUMIR!

Why does this approach scale? Each node stores only one-hop neighborhood information, namely: the id of its parent in Ctree, the IPv6 address ranges of its children in the IPtree and the IPv6 address ranges of its (temporary) children in the RCtree. The memory footprint at each node is therefore $O(k)$, where k is the number of children at any given moment in time.

Why is this approach robust to network dynamics? Routing is performed using the address hierarchy represented by the IPtree, so whenever a link or node fails, messages addressed to destinations in the corresponding subtree may be lost. Matrix uses the reverse collection tree and the local broadcast mechanism to forward messages to the right subtree, as long as an alternative route exists in the collection tree. Note that this two-hop rerouting mechanism does not guarantee that all messages will be delivered. Nevertheless, we argue that the probability that the message will be forwarded to the appropriate subtree is high. The intuition behind this argument is that, given that the network is wireless and, therefore, the number of independent neighbors of a node is limited by a small constant TODO-REF, chances are that, given a failed node or link, its parent in the IPtree will be a one-hop neighbor of the alternative parent replacing it in the collection tree. We show through simulations that this intuition is, in fact, correct, and MATRIX achieves $\geq 95\%$ TODO-CHECK end-to-end packet delivery in wireless low power networks with increasing failure rate and failure duration. TODO: RESUMIR!

We evaluated the proposed protocol both analytically and by simulation. Even though Matrix is platform-independent, we implemented it as a subroutine of CTP on TinyOS and conducted scalability and stress tests. Our analysis showed that, when it comes to any-to-any communication, MATRIX presents significant gains in terms of reliability (higher any-to-any message delivery) and scalability (presenting a constant, as opposed to linear, memory complexity at each node), at a low cost of additional control messages (needed to build IPtree and maintain RCtree), when compared to other available protocols, such as XCTP TODO-CHECK-XCTP and RPL. TODO: include some comparison NUMBERS with RPL/XCTP.

To sum up, Matrix achieves five essential goals that motivated our work:

- **Any-to-any routing:** Matrix enables end-to-end connectivity between hosts located within or outside the 6LoWPAN.
- **Memory efficiency:** Matrix uses compact routing tables and, therefore, is scalable to very large networks.
- **Reliability:** Matrix achieves 100% delivery without end-to-end mechanisms, and delivers $\geq 95\%$ of end-to-end packets when a route exists under challenging network conditions.
- **Communication efficiency:** Matrix uses adaptive beaconing based on Trickle algorithm [11] to minimize the number of control messages in dynamic network topologies (except with node mobility).
- **Hardware independence:** Matrix does not rely on specific radio chip features, and only assumes an underlying collection tree structure.

The rest of this paper is organized as follows. In Section 2 we describe the Matrix protocol design. In Section 3, we analyze the message complexity of the protocol. In Section 4 we present our analytical and simulation results. In Section 5 we discuss some related work. Finally, in Section 6 we present the concluding remarks.

2. DESIGN OVERVIEW

The objective of Matrix is to enable any-to-any routing in an underlying data collection protocol for 6LoWPAN, such as CTP and RPL, while preserving memory and message efficiency, as well as adaptability to networks topology dynamics². Matrix is a network layer protocol that works together with a routing protocol. Figure 1 illustrates the protocol's architecture, which is divided into: *routing engine* and *forwarding engine*. The routing engine is responsible for the address space partitioning and distribution, as well as routing tables maintenance. The forwarding engine is responsible for application packet forwarding.

Matrix is comprised of the following execution phases:

1. **Collection tree initialization:** the collection tree (Ctree) is built by the underlying collection protocol; each node achieves a stable knowledge about who its

²Note that Matrix is not designed to address scenarios with node mobility, but only to work with network topology dynamics caused by changes in link quality, as well as node and link failures.




Figure 1: Matrix protocol's architecture.

parent is; adaptive beaconing based on Trickle algorithm [11] is used to define stability;

2. **Descendants convergecast, IPv6 tree broadcast:** once the collection tree is stable, the address hierarchy tree (IPtree) is built using MHCL [16]; this phase also uses adaptive beaconing to handle network dynamics; by the end of this phase, each node has received an IPv6 address range from its parent and each non-leaf node has partitioned its own address space among its children; the resulting address hierarchy is stored in the distributed IPtree, which initially has the same topology as Ctree, but in reverse, top-down, direction.
3. **Standard routing:** bottom-up routing is done using the collection tree, Ctree, and top-down routing is done using the address hierarchy represented by the IPtree; any-to-any routing is performed by combining bottom-up forwarding, until the least common ancestor of sender and receiver, and then top-down forwarding until the destination.
4. **Alternative top-down routing table upkeep:** whenever a node changes its parent in the initial collection tree, it starts sending beacons to its new parent in Ctree, requesting to upkeep an entry in its routing table with its own IPv6 range; such new links in Ctree, in reverse direction, comprise the RCtree routing tables for alternative (top-down) routing;
5. **Alternative top-down routing via local broadcast:** whenever a node fails to forward a data packet to the next hop/subtree in the IPtree, it broadcasts the packet to its one-hop neighborhood; upon receiving a local broadcast, all neighbors check if the destination IPv6 belongs to an address range in their RCtree table; if positive, the packet is forwarded to the current subtree of IPtree, otherwise, the packet is dropped; we give a geometric argument and show through simulations that such events are rare.

Next we describe the architecture of Matrix in more detail.

2.1 IPv6 multihop host configuration

Matrix is built upon the idea of IPv6 hierarchical address allocation, proposed in [16]. Once the collection tree is stable, the address space available to the border router of the 6LoWPAN, for instance the 64 least-significant bits of the IPv6 address (or a compressed 16-bit representation of the latter), is hierarchically partitioned among nodes in the collection tree. The (top-down) address distribution is preceded by a (bottom-up) convergecast phase, in which each node counts the total number of its descendants, i.e., the size of the subtree rooted at itself, and propagates it to its (preferred) parent. Each node saves the number of descendants of each child. Once the root has received the (aggregate) number of descendants of its k children, it partitions the available address space into k ranges of size proportional to the size of the subtree rooted at each child, leaving a portion of the space as reserve for possible latecoming connections (see Figure 2). Each node repeats the address space partitioning procedure upon receiving its own address range from the parent and sends the proportional address ranges to the respective children, until all nodes have received an address. Since the address allocation is performed in a hierarchical way, each entry in the routing table aggregates the addresses of all destination nodes in the subtree rooted at the corresponding child-node. If a new node connects to the tree after the aggregation phase, it receives an address range from the reserve space of the respective parent node (the details of the communication routines used in this phase are described in detail in [16]).




Figure 2: Example of hierarchical address assignment: (simplified) scenario with 8-bit available address space at the root and 6.25% of address reserve for delayed connections at each node.

After the address configuration phase, the network initialization is done. Each node has built the IPtree routing table with the address range of each child (see Figure 3). All table entries are sorted in increasing order of addresses. In this way, message forwarding can be performed in linear time using one comparison operation per table entry.

2.2 Control plane: distributed tree structures

After the network is initialized and all nodes have received an IPv6 address range, three simultaneous distributed trees are maintained on all nodes in the 6LoWPAN:

- Ctree: the collection tree, maintained by the underlying collection protocol (CTP/RPL).

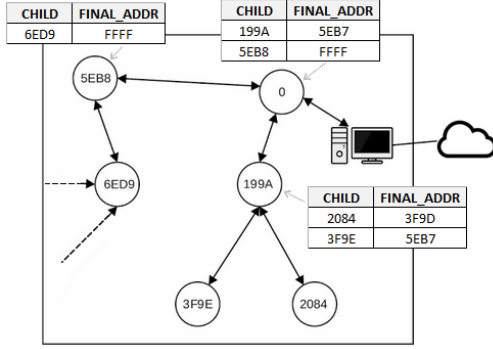


Figure 3: Example of routing tables in IPtree.

- IPtree: the IPv6 address tree, built during the network initialization phase and kept static afterwards, except when new nodes join the network, in which case they receive an IPv6 range from the reserve space of the respective parent node in the collection tree.
- RCTree: the reverse collection tree, reflecting the dynamics of the collection tree in the reverse direction.

Initially, IPtree has the same topology as the reverse-collection tree $Ctree^R$, and RCTree has no links (see Figure 4(a) and 4(b)).

$$IPtree = Ctree^R \text{ and } RCTree = \emptyset$$

Whenever a change occurs in one of the links in Ctree, the new link is added in reverse direction into RCTree and maintained as long as this topology change persists (see Figures 4(c) and 4(d)).

$$RCTree = Ctree^R \setminus IPtree$$

Therefore RCTree is not really a tree, since it contains only the reversed links present in Ctree but not in IPtree. Nevertheless its union with the “working” links in IPtree is, in fact, a tree, which is used in the alternative top-down routing:

$$RCTree \cup (IPtree \cap Ctree^R) : \text{alternative routing tree.}$$

Each node n_i maintains the following information:

- $CTparent_i$: the ID of the current parent in the dynamic collection tree;
- $IParent_i$: the ID of the node that assigned n_i its IPv6 range initially $CTparent_i = IParent_i$;
- $IPchildren_i$: the *standard* (top-down) routing table, with address ranges of each one-hop descendent of n_i in the IPtree, see Figure 3);
- $RChildren_i$: the *alternative* (top-down) routing table, with address ranges of one-hop descendants in the RCTree.

Note that, each node stores only one-hop neighborhood information, so the memory footprint is $O(k)$, where k is the number of a node’s children at any given moment in time, which is optimal.




Figure 4: RCTree example: before and after two links change in the collection tree.

The routing engine (see Figure 1) is responsible for creating and maintaining the IPtree and RCTree routing tables. IPtree is created during the network initialization phase, while RCTree is updated dynamically to reflect changes in the network’s link qualities. Whenever a node n_i has its $CTparent_i$ updated, and the current parent is different from its $IParent_i$ ($IParent_i \neq CTparent_i$), n_i starts sending periodic beacons to its new parent, with regular interval δ , which is the maximum time interval of the Trickle timer (we used the same value provided by CTP/RPL). Upon receiving a beacon (from a new child in the collection tree), a node ($n_j = CTparent_i$) creates and keeps an entry in its alternative routing table $RChildren_j$ with the IPv6 address range of the subtree of n_i . As soon as n_i stops using n_j as the preferred parent, it stops sending beacons to n_j . If no beacon is received from n_i after $2 \times \delta$ ms, this routing entry is deleted. Therefore, links in RCTree are temporary and are deleted when not present in neither the collection nor the IP trees.

2.3 Data plane: any-to-any routing

The forwarding engine (see Figure 1) is responsible for application packet forwarding. Any-to-any routing is performed by combining bottom-up forwarding, until the least common ancestor of sender and receiver, and then top-down forwarding until the destination. Upon receiving an application layer packet, each node n_i verifies whether the destination IPv6 address falls within some range $j \in IPchildren_i$: if yes, then the packet is forwarded (downwards) to node n_j , otherwise, the packet is forwarded (upwards) to $CTparent_i$. Note that, since each node has an IPv6 address, in contrast to collection protocols, such as CTP and RPL, in Matrix, every node can act as a destination of messages originated inside and outside of the 6LoWPAN.

Each forwarded packet requests an acknowledgment from the next hop and can be retransmitted up to 30 times (similarly to what is done in CTP [6]). If thereafter no acknowledgment is received, then the node performs a *local broadcast*, looking for an alternative next hop in the RCTree table of a (one-hop) neighbor. The *alternative routing* process is described in detail below.

2.4 Fault tolerance and network dynamics

So why is Matrix robust to network dynamics? Note that, since routing is based on the hierarchical address allocation, if a node with the routing entries necessary to locate the next subtree becomes unreachable for longer than approximately one second (failures that last less than 1 s are effectively dealt with by retransmission mechanisms available in standard routing protocols), messages with destinations in that subtree are dropped.

When a node or link fails or changes in Ctree, RCtree reflects this change, and packets are forwarded from IPtree to RCtree via a local broadcast. The node that receives a local broadcast checks in its RCtree whether it knows the subtree of the destination IPv6 address: if yes, then it forwards the packet to the right subtree, and the packet continues its path in the IPtree until the final destination.




Figure 5: Alternative top-down routing.

Consider the following scenario: node X receives a packet with destination IPv6 address D (see Figure 5(a)). After consulting its standard routing table $IPchildren_X$, X forwards the packet to C. However, the link $X \Rightarrow C$ fails, for some reason, and C does not reply with an acknowledgement. Then, X makes a constant number (e.g., 30 times in CTP) of retransmission attempts. Meanwhile, since node C also lost its connection to X, it decides to change its parent in the collection tree to node A (see Figure 5(b)). Having changed its parent, C starts sending beacons to A, which creates an entry in its alternative routing table $RCchildren_A$ for the subtree rooted at C, and keeps it as long as it receives periodic beacons from C (which will be done as long as $CTparent_C = A$).

Having received no ack from C, X activates the *local broadcast* mode: it sets the message's type to "LB" and broadcasts it to all its one-hop neighbors (see Figure 5(c)). Upon re-

ceiving the local broadcast, node A consults its alternative routing table and finds out that the destination address D falls within the IPv6 address range C. It then forwards the packet to C, from where the packets follow along its standard route in the subtree of C (see Figure 5(d)).

Note that this mechanism does not guarantee that the message will be delivered. If no one-hop neighbor of X had the address range of C in its alternative routing table, then the packet would be lost. Nevertheless, we argue that the probability that the message will be forwarded to the appropriate subtree is high.

2.5 Alternative routing: geometric rationale

The success of the local broadcast mechanism lies in the ability to forward messages top down along the IPtree, in spite of one or more link or node failures on the way. Matrix is designed to handle (non-adjacent) link or node failures and relies on a single local broadcast and temporary reverse collection links (RCtree).

Consider once again the scenario illustrated in Figure 5. When a node X is unable to forward a packet to the next hop, it activates the local broadcast mechanism, and it becomes essential that one of X's one-hop neighbors (in this case A) has replaced X as a parent of C in the collection tree. Therefore, given that the new parent of C is A, it becomes essential that X and A are neighbors. We argue that it is unlikely that this is not the case.




Figure 6: UDG model: the number of independent neighbors of X is at most 5.

Our argument is of geometric nature. Since the considered 6LoWPAN is wireless, we show our argument in a unit disk graph (UDG) model [2]. We use the fact that the number of independent neighbors of any node in a UDG is bounded by a small constant, namely 5. The proof of this fact is sketched in Figure 6: consider a node X and its neighbor A. Any node located inside the grey region is a neighbor of both X and A, so any neighbor of X that is independent of (not adjacent to) A has to be outside the grey area and inside the circle around X. Let's call this neighbor B. The next independent neighbor of X has to be located outside the 60 degree sector that starts at B, and so on. This procedure can be repeated no more than 5 times, before the 360 degrees around X are covered.

Given that the maximum number of neighbors that do not know each other is very small, for any possible node distribution and density around X, the probability that two neighbors of X are independent is low. In Figure 5(c), since

both X and A are neighbors of C, the probability that they are themselves neighbors is high. Similar arguments can be used to back the effectiveness of the local broadcast mechanism when dealing with different non-adjacent link and node failures.

Note that this reasoning is only valid in an open space without obstacles and, even then, does not guarantee that the message will be delivered. Nevertheless, our experiments show that this intuition is in fact correct, and Matrix has a TODO-CHECK! 100% message delivery success in scenarios with node failures of increasing frequency and duration.

3. COMPLEXITY ANALYSIS

In this section we assume a synchronous communication model with point-to-point message passing. In this model, all nodes start executing the algorithm simultaneously, and time is divided into synchronous rounds, i.e., when a message is sent from node v to its neighbor u at time-slot t , it must arrive at u before time-slot $t + 1$.

We first analyze the message and time complexity of the IPv6 address allocation phase of Matrix, and then look into the message complexity of the control plane of Matrix after the network initialization phase.

Note that Matrix requires that an underlying acyclic topology (Ctree) has been constructed by the network before the address allocation starts, i.e., every node knows who its parent in the Ctree is. Moreover, one of the building blocks of Matrix is the IPv6 multihop host configuration, performed by MHCL [15].

THEOREM 1. [15] *For any network of size n with a spanning collection tree Ctree rooted at node root, the message and time complexity of Matrix protocol in the address allocation phase is $Msg(Matrix^{IP}(Ctree, root)) = O(n)$ and $Time(Matrix^{IP}(T, root)) = O(depth(Ctree))$, respectively. This message and time complexity is asymptotically optimal.*

PROOF. The address allocation phase is comprised of a tree broadcast and a tree convergecast. In the broadcast operation, a message (with address allocation information) must be sent to every node by the respective parent, which needs $\Omega(n)$ messages. Moreover the message sent by the root must reach every node at distance $depth(Ctree)$ hops away, which needs $\Omega(depth(Ctree))$ time-slots. Similarly, in the convergecast operation, every node must send a message to its parent after having received a message from its children, which needs $\Omega(n)$ messages. Also, a message sent by every leaf node must reach the root, at distance $\leq depth(Ctree)$, which needs $\Omega(depth(Ctree))$ time-slots. \square

Next, we examine the communication cost of the routines involved in the alternative routing, performed in the presence of persistent node and link failures.

THEOREM 2. *Consider a network with n nodes and a failure event that causes \mathcal{L}_{CT} links to change in the collection tree Ctree for at most Δ ms. Moreover, consider a beacon interval of δ ms. The control message complexity of Matrix to perform alternative routing is $Msg(Matrix^{RC}) = O(n)$.*

PROOF. Consider the \mathcal{L}_{CT} link changes in the collection tree Ctree. Note that $\mathcal{L}_{CT} = O(n)$, since Ctree is acyclic and, therefore, has at most $n - 1$ links. Every link that was changed must be inserted in the RCtree table of the

respective (new) parent and kept during the interval Δ using regularly sent beacons from the child to the parent. Given a beacon interval of δ , the total number of control messages is bounded by $\Delta/\delta \times \mathcal{L}_{CT} = O(n)$. \square

Note that, in reality, the assumptions of synchrony and point-to-point message delivery do not hold in a 6LoWPAN. The moment in which each node joins the tree, varies from node to node, such that nodes closer to the root tend to start executing the address allocation protocol earlier than nodes farther away from the root. Moreover, collisions and node and link failures can cause delays and prevent messages from being delivered. We analyze the performance of Matrix in an asynchronous model with collisions and transient node and link failures of variable duration through simulations in Section 4.

4. EVALUATION

In this section we evaluate the performance of Matrix through simulations. In Section 4.1 we describe the simulations setup and used parameters. In Section 4.2 we evaluate the number of beacons transmitted, the usage of routing tables and memory footprint, the downward routing success and compare Matrix with two state-of-the-art protocols: RPL [19] and XCTP [18].

4.1 Simulation setup

Matrix and XCTP were implemented in TinyOS [10] and we run the experiments on the simulator for L2Ns TOSSIM [9]. RPL was implemented in Contiki [4] and was simulated on Cooja [5]. Table 1 presents default simulation parameters, used for each protocol, in a nonfaulty scenario. We use the *LinkLayerModel* tool from TinyOS to generate the topology and connectivity model.

We simulate a range of faulty scenarios, based on experimental data collected from TelosB sensor motes, deployed in an outdoor environment [1]. In each scenario, after every 60 seconds of simulation, each node shutdowns its radio with probability σ and keeps the radio off for a time interval uniformly distributed in $[\varepsilon - 5, \varepsilon + 5]$ seconds (see Table 2). The first scenario (*Scn1*) represents a scenario without failures. The following scenarios represent a combination of the values of σ and ε .

We develop an application in which each node sends 10 messages to the root, and the root acknowledges each of these messages. Nodes start sending application messages 90 seconds after the simulation has started. The entire simulation takes 20 minutes and we run each simulation 10 times. For each graph presented, the curve or bars represents the average and the error bars is the confidence interval of 95%.

4.2 Results

To evaluate the usage of routing table, we compare the amount of entries used by each protocol. Each routing table has the same size: 20 entries. In Figure 7 we compare MATRIX, XCTP, RPL, and MHCL. By MHCL we mean MATRIX in a non-faulty scenario, therefore no alternative table entry is used. MATRIX was simulated using the 10th scenario presented in Table 2. We measure the routing table usage for each node per minute and took the average usage in 20 minutes for each node. Since we run 10 experiments with 100 nodes in each, for each protocol, we have 1000 values.

Table 1: Values for Simulation Parameters

Parameter	Value
Base Station	1 center
Number of Nodes	100
Radio Range (m)	100
Density ($nodes/m^2$)	10
Number of experiments	10
Path Loss Exponent	4.7
Power decay (dB)	55.4
Shadowing Std Dev (dB)	3.2
Simulation duration	20 min
Application messages (node to root + ack)	10 per node

Table 2: Values for Faulty Scenarios

$\sigma \backslash \varepsilon$	10 sec.	20 sec.
1%	Scenario 2	Scenario 3
5%	Scenario 4	Scenario 5
10%	Scenario 6	Scenario 7

Then, we calculated the cumulative distribution function, as presented in Figure 7.

As we can see, in Figure 7, RPL is the only protocol that uses 100% of table entries for some nodes. RPL always installs all possible downward routes, causing intermediate nodes not being able to store all downward routes, thus, leading to packages loss. XCTP stores reverse routes only when required. Because of that, the number of routing entries used by XCTP depends on the number of flows (nodes transmitting data to root), unlike MATRIX and RPL. Since the simulated flows were widely spaced during the simulation time, the XCTP was able to perform efficiently. The difference between MHCL and MATRIX is small and is due to the use of RCTree routing table in case of failure.

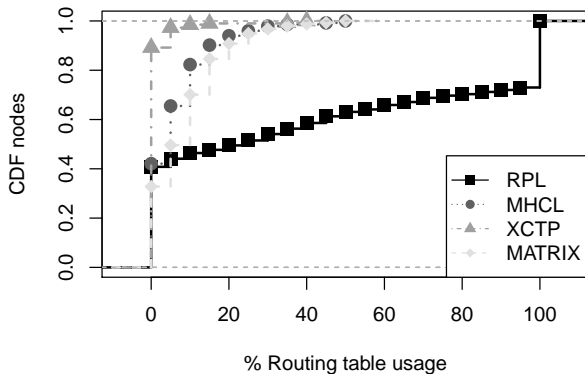


Figure 7: Routing table usage per CDF of nodes for MATRIX, XCTP, MHCL and RPL. For MATRIX we evaluated the 10th scenario.

Figure 8 presents the control traffic from our experiments. The control traffic for all three protocols are high at the network start up, but they stabilize over time, be-

cause RPL and CTP both uses Trickle Timer [11]. MATRIX sends more beacons than XCTP because we create reverse routes for the entire network, while XCTP create it on demand. However, MATRIX sends fewer control packet than RPL because MATRIX only sends additional beacons until all nodes are addressed and in case of failures. RPL presents the highest control overhead because its default parameters to enable downward routes require more beacons to install these rules. **refazer analise apos imagem final**

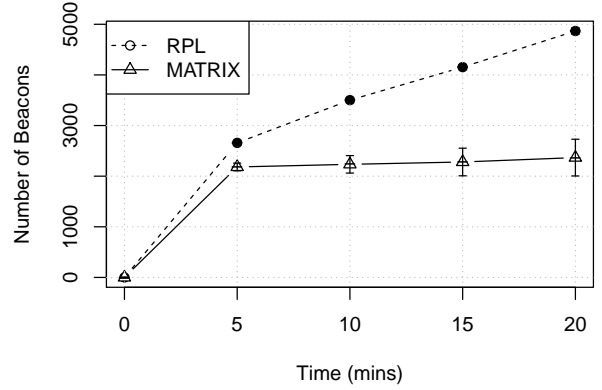


Figure 8: Number of control packets.

Figure 9 presents RAM and ROM footprint sizes of components in the protocol stack. We can see that MATRIX adds only a little more than 7KB of code to CTP, allowing this protocol to perform any-to-any communication with high scalability. When compared with RPL, which has a downward routing mechanism similar to MATRIX, we can see that MATRIX requires smaller amounts of RAM. Our protocol presents RAM requirements close to XCTP while provides scalability, unlike XCTP.




Figure 9: Code and memory footprint in bytes.

Figure 10 presents a comparison between the protocols, for the downward routing success, i. e., the number of application messages sent that were received by the destination.

Fazer a análise de cada protocolo. We call inevitable losses the number of messages that will be lost due to the failure of the destination node, i. e., the destination node could not be accessed even with alternative routes. When this occur, the message will be lost. To infer the number of inevitable losses, we calculated the average time that each of the 99 nodes may be unavailable because of the failure model and how many messages can be targeted to it on that time interval. The remain of messages is lost due to failures on the path or collisions.

avaliar após figura oficial

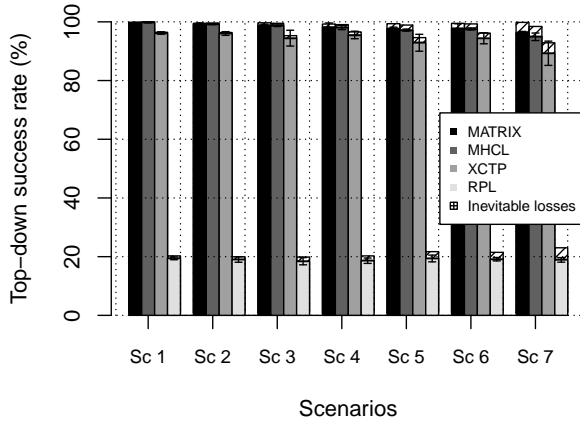


Figure 10: Top-down Success Rate.

5. RELATED WORK

AODV[17] and DSR[7] are traditional wireless protocols that allow any-to-any communication, but they were designed for 802.11 and require large amount of states or apply several overheads on the packet header. In the context of low-power and lossy networks, CTP[6] and CodeDrip[8] were designed for bottom-up and top-down data flows, respectively. They support communication in only one direction.

State-of-the-art routing protocols for 6LoWPAN that enable any-to-any communication are RPL[19], XCTP[18], and Hydro[3]. RPL allows two modes of operation (storing and non-storing) for downwards data flows. The non-storing mode is based on source routing, and the storing mode proactively maintains an entry in the routing table of every node on the path from the root to each destination, which is not scalable to even moderate-size networks. XCTP is an extension of CTP, and is based on a reactive reverse collection route creating between the root and every source node. An entry in the reverse-route table is kept for every data flow at each node on the path between the source and the destination, which is also not scalable in terms of memory footprint. XCTP is suitable for applications that require a small number of reliable end-to-end data flows. Moreover, communication needs to be started by the source nodes. The authors argue that reverse routes can be pro-actively maintained, but additional beacons are required. Hydro protocol, like RPL, is based on a DAG (directed acyclic graph) for

bottom-up communication. Source nodes need to send periodically reports to the border router, which builds a global view (typically incomplete) of the network topology.

Some more recent protocols [14, 13, 12] modified RPL to include new features. In [14], a load-balance technique is applied over nodes to decrease power consumption. In [13, 12], they provide multi-path routing protocols to improve throughput and fault tolerance. Matrix differs from previous work by providing a reliable and scalable solution for any-to-any routing in 6LoWPAN, both in terms of routing table size and control message overhead.

6. CONCLUSIONS

In this paper we propose Matrix: a novel routing protocol that is built upon a data collection structure and is comprised of two phases: (1) network initialization, in which hierarchical IPv6 addresses, which reflect the topology of the underlying wireless network, are assigned to nodes in a multihop way; and (2) reliable any-to-any communication, which enables message and memory-efficient implementation of a wide range of new applications for 6LoWPAN.

7. ADDITIONAL AUTHORS

8. REFERENCES

- [1] N. Baccour, A. Koubâa, L. Mottola, M. A. Zúñiga, H. Youssef, C. A. Boano, and M. Alves. Radio link quality estimation in wireless sensor networks: A survey. *ACM Trans. Sen. Netw.*, 8(4):34:1–34:33, Sept. 2012.
- [2] B. N. Clark, C. J. Colbourn, and D. S. Johnson. Unit disk graphs. *Discrete Math.*, 86(1-3):165–177, Jan. 1991.
- [3] S. Dawson-Haggerty, A. Tavakoli, and D. Culler. Hydro: A hybrid routing protocol for low-power and lossy networks. In *Smart Grid Communications (SmartGridComm)*. IEEE, 2010.
- [4] A. Dunkels, B. Gronvall, and T. Voigt. Contiki - a lightweight and flexible operating system for tiny networked sensors. In *Proceedings of the 29th Annual IEEE International Conference on Local Computer Networks, LCN'14*, pages 455–462, Washington, DC, USA, 2004. IEEE Computer Society.
- [5] J. Eriksson, F. Österlind, N. Finne, N. Tsiftes, A. Dunkels, T. Voigt, R. Sauter, and P. J. Marrón. Cooja/mspsim: Interoperability testing for wireless sensor networks. In *Proceedings of the 2Nd International Conference on Simulation Tools and Techniques, Simutools'09*, pages 27:1–27:7, 2009.
- [6] O. Gnawali, R. Fonseca, K. Jamieson, D. Moss, and P. Levis. Collection tree protocol. In *Proceedings of the 7th ACM Conference on Embedded Networked Sensor Systems, SenSys '09*, pages 1–14, 2009.
- [7] D. Johnson, Y. Hu, and D. Maltz. The dynamic source routing protocol (DSR) for mobile ad hoc networks for IPV4. *RFC: 4728*, 2007.
- [8] N. d. S. R. Júnior, M. A. Vieira, L. F. Vieira, and O. Gnawali. Codedrip: Data dissemination protocol with network coding for wireless sensor networks. In *Wireless Sensor Networks*, pages 34–49. Springer, 2014.

- [9] P. Levis, N. Lee, M. Welsh, and D. Culler. Tossim: Accurate and scalable simulation of entire tinyos applications. In *Proceedings of the 1st International Conference on Embedded Networked Sensor Systems*, SenSys '03, pages 126–137, New York, NY, USA, 2003. ACM.
- [10] P. Levis, S. Madden, J. Polastre, R. Szewczyk, K. Whitehouse, A. Woo, D. Gay, J. Hill, M. Welsh, E. Brewer, et al. Tinyos: An operating system for sensor networks. In *Ambient intelligence*. Springer, 2005.
- [11] P. Levis, N. Patel, D. Culler, and S. Shenker. Trickle: A self-regulating algorithm for code propagation and maintenance in wireless sensor networks. In *Proceedings of the 1st Conference on Symposium on Networked Systems Design and Implementation - Volume 1*, NSDI'04, pages 2–2, 2004.
- [12] M. A. Lodhi, A. Rehman, M. M. Khan, and F. B. Hussain. Multiple path rpl for low power lossy networks. In *Wireless and Mobile (APWiMob), 2015 IEEE Asia Pacific Conference on*, pages 279–284, Aug 2015.
- [13] M. N. Moghadam, H. Taheri, and M. Karrari. Multi-class multipath routing protocol for low power wireless networks with heuristic optimal load distribution. *Wirel. Pers. Commun.*, 82(2):861–881, May 2015.
- [14] U. Palani, V. Alamelumangai, and A. Nachiappan. Hybrid routing and load balancing protocol for wireless sensor network. *Wireless Networks*, pages 1–8, 2015.
- [15] B. Peres and O. Goussevskaia. MHCL: IPv6 Multihop Host Configuration for Low-Power Wireless Networks . Technical report, Federal University of Minas Gerais, Department of Computer Science, 2016.
- [16] B. S. Peres and O. Goussevskaia. Ipv6 multihop host configuration for low-power wireless networks. In *Computer Networks and Distributed Systems (SBRC), 2015 Brazilian Symposium on*, May 2015.
- [17] C. Perkins, E. Belding-Royer, and S. Das. Ad hoc on demand distance vector (AODV) routing (RFC 3561). *IETF MANET Working Group*, 2003.
- [18] B. P. Santos, M. A. Vieira, and L. F. Vieira. extend collection tree protocol. In *Wireless Communications and Networking Conference (WCNC), 2015 IEEE*, pages 1512–1517, March 2015.
- [19] T. Winter, P. Thubert, A. Brandt, J. Hui, R. Kelsey, P. Levis, K. Pister, R. Struik, J. Vasseur, and R. Alexander. RPL: IPv6 Routing Protocol for Low-Power and Lossy Networks. RFC 6550 (Proposed Standard), 2012.